

A Short Practical Introduction to Bayesian Analysis

Opher Donchin

24–25 September 2026

A short worked example

The Bayesian toolkit

The Bayesian workflow

Working through the workflow: golf putting

Preparing for active learning: sleep deprivation


A short worked example

Dose response

Four doses, twenty animals

A toxin is given to rats at four doses, five rats per dose.

log dose	rats	deaths
-0.86	5	0
-0.30	5	1
-0.05	5	3
0.73	5	5




Data: Racine-Poon et al. (1986); Gelman & Vehtari, *Bayesian Workflow*, §3.5

Four doses, twenty animals

A toxin is given to rats at four doses, five rats per dose.

log dose	rats	deaths
-0.86	5	0
-0.30	5	1
-0.05	5	3
0.73	5	5



Data: Racine-Poon et al. (1986); Gelman & Vehtari, *Bayesian Workflow*, §3.5

A generative model for mortality

$$y_j \sim \text{Binomial}(n_j, p_j)$$

$$\text{logit}(p_j) = \alpha + \beta x_j$$

$$\alpha \sim \text{Normal}(0, 5)$$

$$\beta \sim \text{HalfNormal}(5)$$


A generative model for mortality

$$y_j \sim \text{Binomial}(n_j, p_j)$$

$$\text{logit}(p_j) = \alpha + \beta x_j$$

$$\alpha \sim \text{Normal}(0, 5)$$

$$\beta \sim \text{HalfNormal}(5)$$




A generative model for mortality

$$y_j \sim \text{Binomial}(n_j, p_j)$$

$$\text{logit}(p_j) = \alpha + \beta x_j$$


$$\alpha \sim \text{Normal}(0, 5)$$

$$\beta \sim \text{HalfNormal}(5)$$




```
with pm.Model() as model:  
    alpha = pm.Normal("alpha", mu=0, sigma=5)  
    beta = pm.HalfNormal("beta", sigma=5)  
    logit_p = alpha + beta * dose  
    pm.Binomial("deaths", n=n, logit_p=logit_p, observed=deaths)
```

Priors: what could we have seen?



Prior predictive data should not


Priors: what could we have seen?



Prior predictive data should not

- violate common sense


Priors: what could we have seen?



Prior predictive data should not

- violate common sense
- make our data highly unlikely

Priors: what could we have seen?



Prior predictive data should not

- violate common sense
- make our data highly unlikely
- confine the model to our data

Sampling and diagnostics


Fitting the model means drawing samples from the posterior distribution.

```
with model:  
    idata = pm.sample(draws=1000, chains=4)
```

Sampling and diagnostics

Fitting the model means drawing samples from the posterior distribution.

```
with model:  
  idata = pm.sample(draws=1000, chains=4)
```




Check before interpreting:

- 0 divergences
- $\hat{R} = 1.00$
- ESS > 1,300 of 4,000 draws

Posterior: which dose-response curves remain plausible?


The bands show uncertainty about the mean number of deaths.




Posterior: which dose-response curves remain plausible?

The bands show uncertainty about the mean number of deaths.

Prior




Posterior




LD50: the quantity we actually care about

- Any quantity can be computed from the parameters:
 - $LD50 = -\alpha/\beta$
 - Once for every posterior draw.



LD50: the quantity we actually care about


- Any quantity can be computed from the parameters:
 - $LD50 = -\alpha/\beta$
 - Once for every posterior draw.
- So every quantity comes with its own uncertainty.




mean 930 mg/ml
90% HDI 716–1,119 mg/ml

LD50: the quantity we actually care about

- Any quantity can be computed from the parameters:
 - $LD50 = -\alpha/\beta$
 - Once for every posterior draw.
- So every quantity comes with its own uncertainty.
- Here *how well* we know LD50 matters as much as its value.




Posterior predictive: can the model reproduce the experiment?



This is what the model predicts we could have seen, *after* fitting.


Posterior predictive: can the model reproduce the experiment?



This is what the model predicts we could have seen, *after* fitting.

- A workflow check meant to be done before looking at the posterior.


Posterior predictive: can the model reproduce the experiment?



This is what the model predicts we could have seen, *after* fitting.

- A workflow check meant to be done before looking at the posterior.
- Workflow is not lockstep: going back to check is always fine.

Posterior predictive: can the model reproduce the experiment?



This is what the model predicts we could have seen, *after* fitting.

- A workflow check meant to be done before looking at the posterior.
- Workflow is not lockstep: going back to check is always fine.
- We understand data better than we understand parameters.

The Bayesian toolkit

Four libraries, one workflow



Write the model, sample the posterior

pymc.io



Choose and check priors

preliz.readthedocs.io



Diagnose, summarize, and plot

python.arviz.org



Regression models from a formula


bambinos.github.io/bambi

Bambi: the same model as one formula


```
priors = {  
    "Intercept": bmb.Prior("Normal", mu=0, sigma=5),  
    "dose": bmb.Prior("HalfNormal", sigma=5),  
}  
model = bmb.Model("p(deaths, n) ~ dose", data,  
                  family="binomial", priors=priors)  
idata = model.fit()  
bmb.interpret.plot_predictions(model, idata, "dose")
```

Bambi: the same model as one formula

```
priors = {  
    "Intercept": bmb.Prior("Normal", mu=0, sigma=5),  
    "dose": bmb.Prior("HalfNormal", sigma=5),  
}  
model = bmb.Model("p(deaths, n) ~ dose", data,  
                  family="binomial", priors=priors)  
idata = model.fit()  
bmb.interpret.plot_predictions(model, idata, "dose")
```



Coding assistants can learn current PyMC practice




github.com/pymc-labs/pymc-modeling


Skill files from PyMC Labs that teach AI coding assistants current PyMC and ArviZ practice.

We used them to build this course's notebooks.


Where to ask, hire, and listen



Ask: PyMC Discourse
discourse.pymc.io



Hire: PyMC Labs
pymc-labs.com



Listen: Learning Bayesian Statistics
learnbayesstats.com


Books worth reading next




McElreath
Statistical Rethinking




Gelman, Hill & Vehtari
Regression and Other Stories



Gelman & Vehtari
Active Statistics



Martin
Bayesian Analysis with Python



Gelman, Vehtari et al.
Bayesian Workflow

The Bayesian workflow

The workflow starts with a four-step loop



Data

The workflow starts with a four-step loop



The workflow starts with a four-step loop



The workflow starts with a four-step loop



The workflow starts with a four-step loop



Even at this level, we allow for


different interpretations · iterative experimental design

Checks make each transition accountable



Validation belongs at each step


Checks make each transition accountable



Validation belongs at each step
Prior predictive

tests model coherence

Checks make each transition accountable



Validation belongs at each step


Prior predictive

Diagnostics

tests model coherence

test whether fitting converged

Checks make each transition accountable



Validation belongs at each step

Prior predictive

Diagnostics


Posterior predictive

tests model coherence

test whether fitting converged

tests fit to the observed data


A failed check tells us where to return



A failed check

means something must change


A failed check tells us where to return



A failed check
Prior predictive

means something must change
revisit assumptions about the model or data


A failed check tells us where to return



A failed check
Prior predictive
Diagnostics

means something must change
revisit assumptions about the model or data
revisit the model or the computation


A failed check tells us where to return



A failed check
Prior predictive
Diagnostics
Posterior predictive

means something must change
revisit assumptions about the model or data
revisit the model or the computation
reassess how the model represents the data

Model comparison asks which useful story predicts best



Leave-one-out cross-validation			
Model	LOO	SE	Δ LOO
Model 1	-124.8	4.5	13.0
Model 2	-112.4	0.3	0.6
Model 3	-111.8	–	0.0

illustrative values; larger LOO is better

The workflow can lead to several **plausible models** with different interpretations. It is also a useful way to think about frequentist “null models.”

Model comparison is an essential tool in Bayesian data analysis.
All models are wrong; some models are informative.


Working through the workflow: golf putting

Two putting datasets let us fit, then genuinely predict



AI-generated course illustration

Berry (1996): professional putting summaries at rounded distances from 2 to 20 feet.





Two putting datasets let us fit, then genuinely predict



AI-generated course illustration

Berry (1996): professional putting summaries at rounded distances from 2 to 20 feet.


Broadie (2018): a later, much larger dataset extending to 75 feet.



Start simple: logistic regression shows what remains unexplained

$$y_j \sim \text{Binomial}(n_j, p_j)$$
$$\text{logit}(p_j) = \alpha + \beta x_j$$

```
with pm.Model() as model:
    alpha = pm.Normal("alpha", 0, 3)
    beta = pm.Normal("beta", 0, 0.5)
    logit_p = alpha + beta *
distance
    pm.Binomial(
        "made", n=attempts,
        logit_p=logit_p,
        observed=made,
    )
```



Start simple. It is usually best to begin with the smallest useful model.

Prior and posterior predictive checks develop intuition for what it can generate.

What it does and does not capture tells us how much remains to explain.

Geometry turns distance into a success probability




$$\theta(x) = \arcsin\left(\frac{R-r}{x}\right)$$
$$p(x) = 2\Phi\left(\frac{\theta(x)}{\sigma}\right) - 1$$

- A highly simplified model of putting physics
- The only parameter is noise in the aiming angle, σ
- These counts overwhelm almost any reasonable prior
- It is still healthy to ask what a reasonable prior implies


Redrawn from Gelman et al., *Bayesian Workflow*, Fig. 25.3

One angular-noise parameter fits the original data remarkably well

```
with pm.Model() as model:
    sigma_deg = pm.LogNormal(
        "sigma_deg", log(2), 0.7
    )
    theta = pm.math.arcsin(
        (cup_radius - ball_radius)
        / distance
    )
    p = 2 * pm.math.invprobit(
        theta / deg2rad(sigma_deg)
    ) - 1
    pm.Binomial(
        "made", n=attempts,
        p=p, observed=made,
    )
```



A trusted model earns a harder test on new data



When we have faith in a model, we can test it against new data as those data become available.

Data: Broadie (2018); comparison follows *Bayesian Workflow*, §25.3

Preparing for active learning: sleep deprivation
